

“Can a Neural Network write poetry? Leveraging next sentence prediction with Transformer GPT-2”
Research Computing Center Workshop — Summer 2020
Jeffrey Tharsen tharsen@uchicago.edu

This Guide + source files and sample outputs for Blake and Shakespeare :

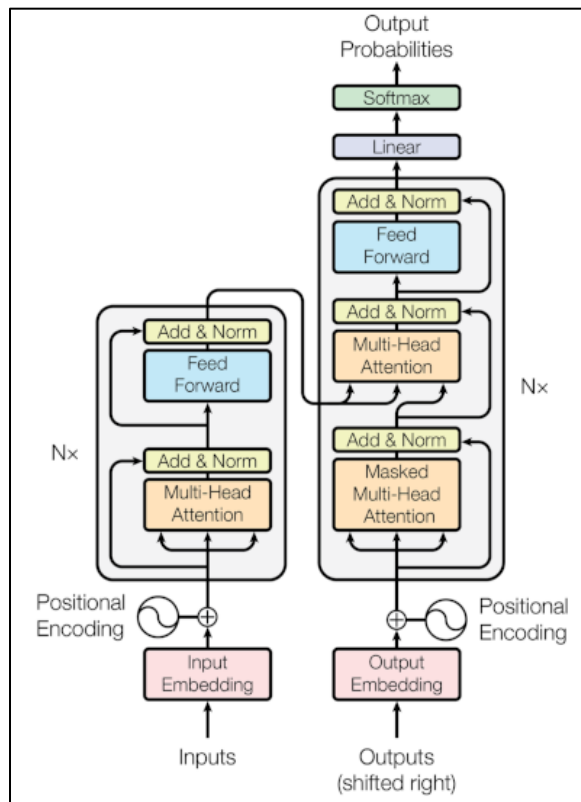
https://github.com/rcc-uchicago/BERT+GPT2_tutorial_Summer2020

Materials from the previous workshop on BERT & GPT-2 :

<https://github.com/rcc-uchicago/BERT-GPT2-Workshop>

1. (For those who haven't seen this yet.)
“This Grad Student Used a Neural Network to Write His Papers” (GPT-2)
<https://futurism.com/grad-student-neural-network-write-papers>
2. **Why GPT-2? What can Trained Transformers achieve? (Encoder-Decoder stacks)**

<https://medium.com/huggingface/encoder-decoders-in-transformers-a-hybrid-pre-trained-architecture-for-seq2seq-af4d7bf14bb8>



<https://openai.com/blog/better-language-models/>

Original source for this tutorial:

<https://medium.com/@ngwaifoong92/beginners-guide-to-retrain-gpt-2-117m-to-generate-custom-text-content-8bb5363d8b7f>

3. Training GPT-2 on your Source Corpus and Producing NN-generated Text

(With Tensorflow – we'll be using Tf 1.13.1 for this tutorial.)

A. Get the GPT-2 framework – from the command line, do:

git clone <https://github.com/nshepperd/gpt-2.git>

(Or go to this link, click on “Code” and download the zip file:

<https://github.com/nshepperd/gpt-2>)

B. `cd gpt-2`

C. Get the model files:

python download_model.py 117M

D. `pip install fire==0.1.3` *(if using Midway, include the --user flag)*

`pip install regex==2017.4.5`

`pip install requests==2.21.0`

`pip install tqdm==4.31.1`

`pip install toposort`

AND

`pip install tensorflow==1.13.1`

*(*** this tutorial does not work with Tensorflow 1.14+ or 2.x ***)*

E. **Put your Source Text File in /gpt-2/src**

e.g. “blake_poems_gpt2.txt”

I split the file into sections with “<|endoftext|>”.

F. Copy the Python code into /src:

`cp encode.py src`

`cp train.py src`

`cp -r models src`

G. `cd src`

If using Midway:

`sinteractive --partition=gpu2 --gres=gpu:1`

`module load Anaconda3/2018.12`

`source activate tf-gpu-1.13.1`

H. `python encode.py [source.txt] [output.npz]`

I. `python train.py --dataset [output.npz]`

It will display samples as it runs; let it run until your loss/avg is under 0.2 .

Train.py will run until you hit Ctrl-C.

- J. Set up the new model directory (I use “blake” here as the name):

```
mkdir models/blake
```

```
cd checkpoint/run1/* models/blake
```

```
cp model-9000* ../../models/blake
```

```
cp checkpoint ../../models/blake
```

```
cp models/117M/encoder.json models/blake
```

```
cp models/117M/hparams.json models/blake
```

```
cp models/117M/vocab.bpe models/blake
```

- K. Now your model is ready to run:

```
python generate_unconditional_samples.py --model_name blake ( > outfile.txt )
```

If you get a model error, edit the /blake/checkpoint file to match the model checkpoint number you wish to use.

- L. Or run (setting **temperature** and **top_k**) :

```
python generate_unconditional_samples.py --temperature 0.8 --top_k 40 --  
model_name blake
```

- M. Finally, you can run interactively, supplying a prompt:

```
python interactive_conditional_samples.py --temperature 0.8 --top_k 40 --  
model_name blake
```